

## DS-2: Heart-Disease Prediction

### Task vii – Automated Feature Extraction & Classification

Dataset DS2.csv

Goal: Predict target (1 = disease, 0 = healthy)

Feature-extraction techniques

1. **PCA** – linear, variance-preserving
2. **t-SNE** – non-linear, manifold-learning (mainly for visual insight, but we'll still feed its low-dim representations to classifiers)

Classifiers

- Logistic Regression (baseline linear)
- Random Forest (non-linear, ensemble)

We'll report **accuracy, precision, recall, F1, confusion-matrix** for each (technique × model) combination.

```
# 1 Environment & libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, confusion_matrix, classification_report)

import matplotlib.pyplot as plt
import seaborn as sns # optional, for prettier matrices
RANDOM_STATE = 42
```

### 2 Load the data

```
df = pd.read_csv('/content/DS2.csv') # adjust path if needed
display(df.head())
print(f'Shape: {df.shape}')
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63.0	1.0	1.0	145.0	233.0	1.0	2.0	150.0	0.0	2.3	3.0	0.0	6.0	0
1	67.0	1.0	4.0	160.0	286.0	0.0	2.0	108.0	1.0	1.5	2.0	3.0	3.0	2
2	67.0	1.0	4.0	120.0	229.0	0.0	2.0	129.0	1.0	2.6	2.0	2.0	7.0	1
3	37.0	1.0	3.0	130.0	250.0	0.0	0.0	187.0	0.0	3.5	3.0	0.0	3.0	0
4	41.0	0.0	2.0	130.0	204.0	0.0	2.0	172.0	0.0	1.4	1.0	0.0	3.0	0

Shape: (303, 14)

### 2.1 Quick sanity checks

```
print(df.info())
print(df.isna().sum())
print(df['target'].value_counts())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    age         303 non-null    float64
1    sex         303 non-null    float64
2    cp          303 non-null    float64
3    trestbps    303 non-null    float64
4    chol        303 non-null    float64
5    fbs         303 non-null    float64
6    restecg     303 non-null    float64
7    thalach     303 non-null    float64
```

```

8   exang      303 non-null    float64
9   oldpeak    303 non-null    float64
10  slope      303 non-null    float64
11  ca         303 non-null    object
12  thal       303 non-null    object
13  target     303 non-null    int64
dtypes: float64(11), int64(1), object(2)
memory usage: 33.3+ KB
None
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
target
0    164
1     55
2     36
3     35
4     13
Name: count, dtype: int64

```

### 3 Pre-processing pipeline

- **Numeric** → StandardScaler
  - **Categorical** → OneHotEncoder
- The target column is target ; all others are features.

```

target_col = 'target'
X = df.drop(columns=[target_col])
y = df[target_col]

numeric_cols = X.select_dtypes(include=['int64', 'float64']).columns.tolist()
categorical_cols = X.select_dtypes(include=['object', 'category']).columns.tolist()

preprocess = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_cols),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_cols)
    ]
)

X.head(20)

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	63.0	1.0	1.0	145.0	233.0	1.0	2.0	150.0	0.0	2.3	3.0	0.0	6.0
1	67.0	1.0	4.0	160.0	286.0	0.0	2.0	108.0	1.0	1.5	2.0	3.0	3.0
2	67.0	1.0	4.0	120.0	229.0	0.0	2.0	129.0	1.0	2.6	2.0	2.0	7.0
3	37.0	1.0	3.0	130.0	250.0	0.0	0.0	187.0	0.0	3.5	3.0	0.0	3.0
4	41.0	0.0	2.0	130.0	204.0	0.0	2.0	172.0	0.0	1.4	1.0	0.0	3.0
5	56.0	1.0	2.0	120.0	236.0	0.0	0.0	178.0	0.0	0.8	1.0	0.0	3.0
6	62.0	0.0	4.0	140.0	268.0	0.0	2.0	160.0	0.0	3.6	3.0	2.0	3.0
7	57.0	0.0	4.0	120.0	354.0	0.0	0.0	163.0	1.0	0.6	1.0	0.0	3.0
8	63.0	1.0	4.0	130.0	254.0	0.0	2.0	147.0	0.0	1.4	2.0	1.0	7.0
9	53.0	1.0	4.0	140.0	203.0	1.0	2.0	155.0	1.0	3.1	3.0	0.0	7.0
10	57.0	1.0	4.0	140.0	192.0	0.0	0.0	148.0	0.0	0.4	2.0	0.0	6.0
11	56.0	0.0	2.0	140.0	294.0	0.0	2.0	153.0	0.0	1.3	2.0	0.0	3.0
12	56.0	1.0	3.0	130.0	256.0	1.0	2.0	142.0	1.0	0.6	2.0	1.0	6.0
13	44.0	1.0	2.0	120.0	263.0	0.0	0.0	173.0	0.0	0.0	1.0	0.0	7.0
14	52.0	1.0	3.0	172.0	199.0	1.0	0.0	162.0	0.0	0.5	1.0	0.0	7.0
15	57.0	1.0	3.0	150.0	168.0	0.0	0.0	174.0	0.0	1.6	1.0	0.0	3.0
16	48.0	1.0	2.0	110.0	229.0	0.0	0.0	168.0	0.0	1.0	3.0	0.0	7.0
17	54.0	1.0	4.0	140.0	239.0	0.0	0.0	160.0	0.0	1.2	1.0	0.0	3.0
18	48.0	0.0	3.0	130.0	275.0	0.0	0.0	139.0	0.0	0.2	1.0	0.0	3.0
19	49.0	1.0	2.0	130.0	266.0	0.0	0.0	171.0	0.0	0.6	1.0	0.0	3.0

Next steps: [Generate code with X](#) [View recommended plots](#) [New interactive sheet](#)

## 4 Train-test split (80 / 20)

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)
print(f'Train: {X_train.shape}, Test: {X_test.shape}')
```

Train: (242, 13), Test: (61, 13)

## 5 Feature Extraction — PCA

We'll keep enough components to explain **95 % variance** (you can tweak).

```
pca = PCA(n_components=0.95, random_state=RANDOM_STATE)

pca_pipe = Pipeline(steps=[
    ('pre', preprocess),
    ('pca', pca)
])

X_train_pca = pca_pipe.fit_transform(X_train)
X_test_pca = pca_pipe.transform(X_test)

print(f'PCA components chosen: {pca.n_components_}')
```

PCA components chosen: 12

## 6 Classification on PCA features

```
models = {
    'LogReg': LogisticRegression(max_iter=1000, random_state=RANDOM_STATE),
    'RF': RandomForestClassifier(n_estimators=300, random_state=RANDOM_STATE)
}
```

```
def evaluate(name, clf, X_train, X_test):
```

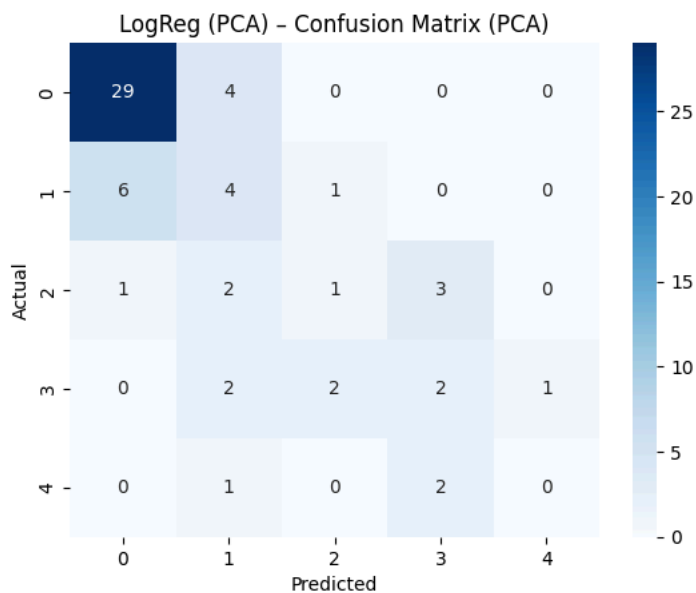
```
clf.fit(X_tr, y_train)
y_pred = clf.predict(X_te)
print(f'\n=== {name} ===')
print(f'Accuracy : {accuracy_score(y_test, y_pred):.4f}')
```

```
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title(f'{name} - Confusion Matrix (PCA)')
plt.xlabel('Predicted'); plt.ylabel('Actual')
plt.show()
```

```
for name, clf in models.items():
    evaluate(f'{name} (PCA)', clf, X_train_pca, X_test_pca)
```



```
=== LogReg (PCA) ===
Accuracy : 0.5902
```



```
=== RF (PCA) ===
Accuracy : 0.5902
```

